
An Audit of Structure-Based HIV Drug-Resistance Triage and Its LLM Explanations

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Antiviral drug resistance forces medicinal chemists to triage candidate inhibitors
2 by how well their binding survives known resistance mutations. We created **Re-**
3 **sistScope**, a target-agnostic pipeline that docks a candidate against wildtype plus
4 a panel of clinical resistance mutants, scores binding robustness ($\Delta\Delta G$), and ex-
5 plains each mutation’s effect with an LLM. Both outputs have ground truth, so we
6 test both on HIV-1 protease (6 PIs) and reverse transcriptase (4 NNRTIs). **The**
7 **docking score is beaten by its own geometry**: ranking mutations by distance
8 to the ligand, same receptor but no docking run, tracks measured fold-resistance
9 better than $\Delta\Delta G$ on both targets. The energy earns its cost only at the top of
10 the ranking in the rigid NNRTI pocket, where each drug’s top ten holds 7–8 real
11 resistance mutations against 3–4 for distance; on the flexible protease site that or-
12 dering reverses. An LLM judge marks 71–72% of the explanations as recovering
13 the correct primary mechanism, and an **attribution ablation** over real, absent and
14 corrupted structural context shows they condition on the pipeline, most strongly
15 where the model’s prior is weakest. A verification using no language model re-
16 produces that ordering. Structural geometry carries the signal; the docking energy
17 mostly does not.

18 1 Introduction

19 Resistance is the central failure mode of antiviral therapy. A medicinal chemist with twenty candi-
20 date inhibitors must decide which few to advance into slow, expensive phenotypic resistance assays,
21 a decision currently made by hand. A natural computational aid is to dock each candidate against
22 a panel of known clinical resistance mutants and rank candidates by how well predicted binding
23 *survives* the panel.

24 Such a tool raises two questions that are usually asked separately, and that both have ground truth:

- 25 1. **Does the prediction hold up?** Stanford HIVdb records *measured* fold-resistance for es-
26 sentially every approved HIV drug [19]. This asks whether the score predicts a measured
27 number, and whether it beats the trivial baselines a chemist could use instead.
- 28 2. **Does the explanation hold up?** If the tool also explains *why* a mutation threatens a com-
29 pound, that explanation can be checked against mechanisms experts have already docu-
30 mented.

31 **Contributions.** (i) An audit of structure-based resistance triage on two HIV targets, showing a
32 docking-free geometric baseline outranks docking $\Delta\Delta G$ on both, with one exception at the top of
33 the ranking in the rigid NNRTI pocket where the energy does earn its cost. (ii) Decision-relevant
34 evaluation in place of a single pooled statistic: per-drug precision@ k , significance-tested baselines,
35 and a dropout sensitivity analysis. (iii) An attribution ablation for the pipeline’s LLM explanations,




Figure 1: **ResistScope pipeline.** A candidate SMILES is docked against wildtype plus a panel of clinical resistance mutants; per-mutation $\Delta\Delta G$ yields a robustness score and, with structural context, an LLM mechanistic explanation. Each output is validated against ground truth: $\Delta\Delta G$ against measured fold-resistance, explanations against documented mechanisms.

transferring the corrupted-context logic of Min et al. [15] to structural input, with a verification that uses no language model. (iv) A reproducible target-agnostic implementation, with all docking results, explanations, ground truth and judge scores released.

2 Related work

Rule-based expert systems are the clinical standard and our strongest baseline. The Stanford HIVdb penalty algorithm [14, 19] and the IAS-USA mutation lists [22] encode decades of genotype-phenotype and clinical-outcome evidence into per-mutation, per-drug scores. They are accurate because that evidence exists; they say nothing about a compound or target for which it does not.

Supervised ML on genotype-phenotype data is the dominant approach and reaches high retrospective accuracy, up to mean AUC ≈ 0.97 with protein language-model embeddings [3, 20, 4, 10]. We do not claim to beat it. It carries one requirement that motivates our work: a labeled resistance dataset *for the drug in question*. Docking is training-free, so it applies to a novel scaffold or a target with no resistance corpus. Whether that applicability is worth anything is what we test.

Structure-based resistance prediction has an established explanation for why protease is hard. The substrate-envelope hypothesis [13, 16] holds that resistance mutations disrupt inhibitor binding where the inhibitor protrudes beyond the consensus substrate volume, while leaving substrate processing intact. That is a coevolutionary, network-level constraint a single-structure pocket-affinity score is not designed to capture, and it predicts our protease result. Receptor flexibility is the standard remedy for rigid-docking failure [2]; we do not employ it, which bounds our negative protease result.

LLM explanation faithfulness. Whether a model’s stated reasoning reflects the computation producing its answer is distinct from whether the answer is right [12, 21]. Our corrupted-context condition follows Min et al. [15], who showed that randomizing in-context labels barely degrades performance, implying the model was not conditioning on them; we apply that test to structural context. LLM-as-judge scores are judge-dependent [24], so we rest our claims on paired within-judge contrasts. Our design is not new; its transfer to *structural* input is, as is a verification that uses no language model at all. Appendix G covers the alternatives to rigid docking and the neighbouring perturb-and-rejudge work.

3 Methods

Pipeline (Fig. 1). A candidate SMILES is embedded and docked (AutoDock Vina [9], or Uni-Dock [23] on GPU) against a cleaned wildtype receptor and a per-drug panel of PDBFixer-built point-mutant receptors; the per-mutation $\Delta\Delta G$ (mutant minus wildtype binding) feeds a 0–100 robustness score. The Vina score is an empirical scoring function, not a binding free energy; we use it as a ranking signal only.

Targets and data. Protease (6 PIs, 1,591 scored pairs) and reverse transcriptase (4 NNRTIs, 1,269 pairs); structures, panel sizes and docking boxes in Appendix D. We scope RT to NNRTIs (allosteric pocket binders): rigid docking is a sound proxy for NNRTI affinity but not for NRTI mechanisms. Measured fold-resistance and DRM labels come from Stanford HIVdb / IAS-USA [19, 22].

Explanations and ground truth. For each mutation we prompt an LLM (Claude Haiku 4.5) with a structural-context record (residue identity, Δ volume/charge/hydrophobicity, ligand distance, direct-contact flag, subpocket) plus the docking $\Delta\Delta G$, for a 2–3 sentence mechanistic hypothesis. Ground truth combines a curated set (IAS-USA + primary literature) with an *agentic* builder that researches each mechanism over live PubMed, cites the papers, and self-verifies grounding. All RT ground truth is agent-built and PR’s is hand-curated, so we lead the faithfulness headline with PR.

Faithfulness and attribution. An LLM-as-judge scores each explanation against the ground-truth mechanism on a 0/1/2 rubric (judge robustness: §4.2). The attribution ablation regenerates explanations under *full* (real context + $\Delta\Delta G$), *minimal* (drug + mutation only), and *corrupted* (pipeline-derived fields deranged across pairs, intrinsic chemistry kept) conditions with the same model and judge; a field-level variant drops one context group at a time.

Predictive benchmark. Every claim carries a permutation p -value or bootstrap 95% CI, scored against docking-free baselines: mutation prevalence, $|\Delta$ volume|, ligand distance, and the Stanford HIVdb penalty. Test protocol, de-confounding, jackknife and the absence of multiplicity correction are in Appendix D.

4 Results

4.1 The predictive benchmark: geometry, not energy

On both enzymes, rigid single-structure docking- $\Delta\Delta G$ is beaten by a baseline that uses the same receptor but never docks anything. The one exception tells us where the energy step is worth its cost.

A docking-free geometric baseline beats $\Delta\Delta G$ globally, on both targets. The field-level ablation of §4.3 finds the *explanations* carried by geometry rather than energy; so is the *ranking*. Scoring each mutation by how close its residue lies to the ligand, with no docking run, raises the magnitude correlation with measured fold-resistance from $\rho = 0.002$ to 0.387 on protease, above HIVdb’s 0.360, and DRM-recovery ROC-AUC from 0.51 [0.46, 0.56] to 0.963 [0.954, 0.971] (RT: 0.70 to 0.800 [0.738, 0.854]). We lead with the magnitude column because DRM labels are enriched for pocket-proximal residues almost by construction, which makes the ROC comparison partly definitional. The ordering is the same on both.

The exception: at the top of the list, on RT, docking wins. ROC-AUC scores a whole ranking, but the tool is read as the top few flagged mutations for one compound, and the two metrics disagree. Taking each drug’s own top 10, docking recovers 7 to 8 genuine clinical DRMs per NNRTI against 3 to 4 for the distance baseline. On protease the result reverses, 0 to 3 for docking against 7 to 10 for distance (Fig. 2; Appendix B). In the NNRTI pocket proximity is necessary but not sufficient, since many close residues are not resistance mutations, and discriminating *within* the pocket is the one thing the energy buys. On protease it buys nothing at any depth of the list.

Why the targets differ (Fig. 4, Appendix A). On protease docking sits at chance (Spearman ≈ 0 ; top-40 enrichment $2.86\times$, $p < 0.001$, so tail enrichment coexists with chance-level global ranking) and de-confounding does not rescue it ($\rho \leq 0$). That is expected: protease resistance is co-evolving and substrate-envelope-constrained [13, 16], violating the locality the method assumes (Appendix E). On RT the same protocol is above chance (enrichment $8.35\times$) and de-confounding *does* rescue the magnitude correlation, $-0.11 \rightarrow \rho = 0.36$ [0.12, 0.56] ($p = 6 \times 10^{-4}$) on single-mutation isolates.

4.2 Explanation faithfulness

Under the Claude-Haiku judge, explanations recover the correct primary mechanism 72% of the time on protease ($n = 46$, hand-curated ground truth) and 71% on RT ($n = 55$, agent-built ground truth, so its absolute value is the weaker of the two); most misses under-specify rather than contradict. The outright misses cluster on distal and inter-subunit DRMs, notably the E138–K101 salt bridge across the p51/p66 interface [8], where the structural context itself is blind rather than the model (Appendix E).

Judge robustness. A second Claude model re-grading all 138 protease ablation items reproduces the §4.3 contrasts, and GPT-5.6 agrees with the Haiku judge 95% within one point but only 48%




Figure 2: **Top of the list, per drug.** Precision@10 for the two structure-only predictors, the depth at which the tool is actually read; Table 4 ranks the whole panel instead. Each line joins one drug’s two values, the tick is its DRM base rate, and * marks $p < 0.05$ by an exact one-sided hypergeometric test against that rate (scripts/15_geometry_baseline.py). The ordering reverses between targets, which is the one place the energy step earns its cost.




Figure 3: **Attribution ablation** (§4.3). (a) The benefit of correct context (full – minimal) by the model’s prior strength: largest where the prior alone is wrong, zero where it is already correct. (b) Field-level ablation: dropping ligand distance costs most and subpocket next (both marginal at $n = 46$); dropping the docking energy or the physicochemistry costs nothing measurable.

124 exactly, Claude being systematically *stricter*. We therefore treat the *absolute* percentage as judge-
 125 dependent [24] and rest our claims on the paired contrasts (Appendix D).

126 **An out-of-family check.** All three judges are LLMs, so the above establishes consistency, not
 127 correctness. Appendix F adds a check involving no language model: whether an explanation’s
 128 claim about ligand contact matches the receptor geometry. It reproduces the ablation ordering (88%
 129 correct under full context, 57% minimal, 55% corrupted; Fisher $p = 5.7 \times 10^{-5}$ full vs. corrupted)
 130 and shows the judge scoring structurally-wrong explanations lower (1.02 vs. 1.51, $p = 6.1 \times 10^{-4}$).

Both results most exposed to the circularity objection therefore survive a measurement taken outside the LLM family.

4.3 Attribution ablation: the pipeline drives faithfulness where the prior is weak

Context is actually used. Pooled across both targets ($n = 101$), full context beats minimal (drug + mutation only) by $\Delta = +0.27$ ($p = 7 \times 10^{-4}$) and beats corrupted context by $\Delta = +0.46$ ($p < 10^{-4}$). Feeding *wrong* structural facts is worse than feeding none, so by the logic of Min et al. [15] the explanations are conditioning on the pipeline rather than pattern-matching the mutation name (Appendix C). *How much is the contact flag?* Deranging geometry *within* contact class, leaving that flag true, recovers most of the loss (1.63 vs. 1.15 corrupted, 1.78 full), but real context still wins ($\Delta = +0.15$, $p = 0.035$). The finer geometry contributes beyond the boolean, though the headline arm overstates it fourfold (Appendix D).

Prior-strength dose-response. The benefit of correct context scales inversely with the model’s prior (Fig. 3a). Where the prior alone is *wrong* (minimal = 0, $n = 13$), context corrects the explanation in 11 of 13 cases (85%); where vague (minimal = 1), 66%; where already correct (minimal = 2), context is at ceiling. Because the raw gain is partly mechanical (a 0 has twice the headroom of a 1) we lead with these ceiling-robust rates; restricting to pairs with room to improve (minimal < 2, $n = 45$), context still helps $\Delta = +0.80$ ($p < 10^{-4}$) and the gradient persists (Spearman -0.40 , $p = 0.006$).

Which fields carry the lift. A field-level ablation (protease, $n = 46$; Fig. 3b) drops one context group at a time. Removing the *ligand distance* costs the most ($\Delta = +0.22$, $p = 0.051$) and the *subpocket* next ($\Delta = +0.15$, $p = 0.09$); removing the *docking* $\Delta\Delta G$ ($\Delta = -0.02$) or the physicochemistry ($\Delta = -0.02$) does nothing measurable. Marginally, and on one target ($n = 46$), that points to geometric localization rather than energy. The $\Delta\Delta G$ null is absence of evidence, and specific to *rigid, single-pose, waters-stripped Vina*: a better energy [2, 11] might contribute where ours does not.

5 Discussion

Both halves point the same way: **structural geometry carries the signal, and the docking energy mostly does not.** Ranking by ligand distance beats docking $\Delta\Delta G$ globally on both targets, and dropping distance or subpocket costs the most faithfulness while dropping $\Delta\Delta G$ costs nothing measurable. Two independent analyses reach one conclusion, a negative result about the method we built. The energy is not worthless. It discriminates *within* the NNRTI pocket, where proximity alone cannot, and that is the whole of what we can show it buys. Structure still beats the non-structural shortcuts where resistance is structurally local; the *energy* step is what fails to earn its cost. The baselines that beat docking need per-drug data a novel compound lacks (Appendix H), and nothing in the $\Delta\Delta G$ path is fitted to resistance labels. Both boundaries follow from the premise rather than surprising it, and so do the individual failures (Appendix E): eight of nine trace to what the pipeline cannot represent rather than to the language model. We offer this from two contrasting HIV enzymes as a hypothesis, not a law: measure the geometry, treat the docking energy as unproven, and distrust structural triage for flexible, co-evolving sites.

Limitations. Docking-energy claims are bounded by this protocol: rigid single-structure Vina, one pose, waters stripped, and 7.1% of RT dockings dropped for returning no pose. The faithfulness half rests on $n = 46$ and $n = 55$ pairs over two enzymes, with load-bearing contrasts marginal ($p = 0.035$, 0.051 , 0.09) and no multiplicity correction. RT ground truth is agent-built, all three judges are language models, and most of the corruption effect is a single contact flag (§4.3). Appendix D details each.

Future work. A *human-expert anchor*, 20 to 30 explanations graded blind on the same rubric by a structural biologist, is the one measurement that would move the faithfulness result from consistent to correct; the sample and primer are released. *More targets* would test geometry-over-energy beyond two HIV enzymes; neuraminidase and SARS-CoV-2 M^{pro} are already wired in. *A better energy* (ensemble docking [2], free-energy calculation [11]) would separate a null about docking from a null about this protocol. *Replication at larger n*, including on RT, would settle the marginal decomposition claims ($p = 0.035$ to 0.09 at $n = 46$); every number here is also retrospective (§4.1).

References

- [1] Afolabi, H., et al. Faithful or just plausible? Evaluating the faithfulness of closed-source LLMs in medical reasoning. *arXiv:2603.13988*, 2026.
- [2] Amaro, R. E., Baron, R., McCammon, J. A. An improved relaxed complex scheme for receptor flexibility in computer-aided drug design. *Journal of Computer-Aided Molecular Design*, 2008.
- [3] Beerenwinkel, N., Däumer, M., Oette, M., Korn, K., Hoffmann, D., Kaiser, R., Lengauer, T., Selbig, J., Walter, H. Geno2pheno: estimating phenotypic drug resistance from HIV-1 genotypes. *Nucleic Acids Research*, 31(13):3850–3855, 2003.
- [4] Blassel, L., Tostevin, A., Villabona-Arenas, C. J., Peeters, M., Hué, S., Gascuel, O. Using machine learning and big data to explore the drug resistance landscape in HIV. *PLOS Computational Biology*, 17(8):e1008873, 2021.
- [5] Passaro, S., et al. Boltz-2: Towards accurate and efficient binding affinity prediction. *bioRxiv*, 2025. doi:10.1101/2025.06.14.659707.
- [6] M. Bran, A., Cox, S., Schilter, O., Baldassari, C., White, A. D., Schwaller, P. Augmenting large language models with chemistry tools. *Nature Machine Intelligence*, 6:525–535, 2024.
- [7] Buttenschoen, M., Morris, G. M., Deane, C. M. PoseBusters: AI-based docking methods fail to generate physically valid poses or generalise to novel sequences. *Chemical Science*, 15(9):3130–3139, 2024.
- [8] Das, K., Arnold, E. HIV-1 reverse transcriptase and antiviral drug resistance. Part 2. *Current Opinion in Virology*, 2013.
- [9] Eberhardt, J., Santos-Martins, D., Tillack, A. F., Forli, S. AutoDock Vina 1.2.0: New docking methods, expanded force field, and Python bindings. *Journal of Chemical Information and Modeling*, 61:3891–3898, 2021. See also Trott, O., Olson, A. J., *Journal of Computational Chemistry*, 31:455–461, 2010.
- [10] Farquhar, H. Protein language model embeddings improve HIV drug resistance prediction: a comprehensive benchmark with attention-based interpretability. *Bioinformatics*, 42(5):btac260, 2026.
- [11] Hauser, K., Negron, C., Albanese, S. K., Ray, S., Steinbrecher, T., Abel, R., Chodera, J. D., Wang, L. Predicting resistance of clinical Abl mutations to targeted kinase inhibitors using alchemical free-energy calculations. *Communications Biology*, 1:70, 2018.
- [12] Jacovi, A., Goldberg, Y. Towards faithfully interpretable NLP systems: How should we define and evaluate faithfulness? In *Proceedings of ACL*, pp. 4198–4205, 2020.
- [13] King, N. M., Prabu-Jeyabalan, M., Nalivaika, E. A., Schiffer, C. A. Combating susceptibility to drug resistance: lessons from HIV-1 protease. *Chemistry & Biology*, 2004.
- [14] Liu, T. F., Shafer, R. W. Web resources for HIV type 1 genotypic-resistance test interpretation. *Clinical Infectious Diseases*, 42(11):1608–1618, 2006.
- [15] Min, S., Lyu, X., Holtzman, A., Artetxe, M., Lewis, M., Hajishirzi, H., Zettlemoyer, L. Rethinking the role of demonstrations: What makes in-context learning work? In *Proceedings of EMNLP*, 2022.
- [16] Nalam, M. N. L., Ali, A., Altman, M. D., Reddy, G. S. K. K., Chellappan, S., Kairys, V., Özen, A., Cao, H., Gilson, M. K., Tidor, B., Rana, T. M., Schiffer, C. A. Evaluating the substrate-envelope hypothesis: structural analysis of novel HIV-1 protease inhibitors designed to be robust against drug resistance. *Journal of Virology*, 84(10):5368–5378, 2010.
- [17] Porbeck, J., Adriano, C. M., Giese, H. From program slices to causal clarity: Evaluating faithful, actionable LLM-generated failure explanations via context partitioning and LLM-as-a-judge. *arXiv:2604.18309*, 2026.

- [18] Ota, R., So, K., Tsuda, M., Higuchi, Y., Yamashita, F. Prediction of HIV drug resistance based on the 3D protein structure: Proposal of molecular field mapping. *PLOS ONE*, 2021. doi:10.1371/journal.pone.0255693.
- [19] Rhee, S.-Y., Taylor, J., Wadhwa, G., Ben-Hur, A., Brutlag, D. L., Shafer, R. W. Genotypic predictors of human immunodeficiency virus type 1 drug resistance. *PNAS*, 103(46):17355–17360, 2006. Stanford HIV Drug Resistance Database, <https://hivdb.stanford.edu>.
- [20] Steiner, M. C., Gibson, K. M., Crandall, K. A. Drug resistance prediction using deep learning techniques on HIV-1 sequence data. *Viruses*, 12(5):560, 2020.
- [21] Turpin, M., Michael, J., Perez, E., Bowman, S. R. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems 36*, 2023.
- [22] Wensing, A. M., Calvez, V., Ceccherini-Silberstein, F., Charpentier, C., Günthard, H. F., Paredes, R., Shafer, R. W., Richman, D. D. 2022 update of the drug resistance mutations in HIV-1. *Topics in Antiviral Medicine*, 30(4):559–574, 2022.
- [23] Yu, Y., Cai, C., Wang, J., Bo, Z., Zhu, Z., Zheng, H. Uni-Dock: GPU-accelerated docking enables ultralarge virtual screening. *Journal of Chemical Theory and Computation*, 2023.
- [24] Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., Zhang, H., Gonzalez, J. E., Stoica, I. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems 36*, 2023.

A Predictive benchmark figures




Figure 4: **Docking $\Delta\Delta G$ across the two targets.** (a) DRM-recovery ROC-AUC with bootstrap 95% CIs: docking- $\Delta\Delta G$ sits on the chance line for protease (0.51) and above it for RT (0.70), but a docking-free ligand-distance baseline beats it on both (Table 5), and the HIVdb expert system wins both. (b) De-confounding: as the fold-resistance target is rebuilt from progressively less co-mutated isolates, RT’s magnitude correlation climbs to $\rho \approx 0.4$ while protease stays ≤ 0 .

B Per-drug breakdown and precision@ k

Table 1 gives the per-drug metrics behind the pooled numbers in §4.1, and Table 2 the decision-relevant precision@ k . Both are produced by `scripts/14_precision_at_k.py`.

Precision@ k ranks a *single* drug’s own mutation panel by predicted $\Delta\Delta G$ (descending, matching the pooled enrichment convention) and reports how many of the top k are known major DRMs. Significance is an exact one-sided hypergeometric test, giving the probability that drawing k mutations at random from that drug’s panel yields at least as many DRMs. That is appropriate at $k = 5$ where a bootstrap CI is uninformative.

The per-drug view shows the pooled result is not an artifact of averaging. Every NNRTI clears its base rate by a wide margin at every k ; among the PIs only atazanavir ($k = 5$) and lopinavir ($k = 20$) reach $p < 0.05$, and indinavir, whose ROC-AUC of 0.388 is *below* chance, returns no DRMs in its top 10. The protease failure is therefore not uniform noise: it is one target on which the method is, for some drugs, actively anti-correlated with the truth.

The last column is the comparison §4.1 rests on, and it runs opposite to the ROC-AUC ordering of Table 5. Distance is the better global ranker on both targets, but at $k = 10$ it wins only on protease (7 to 10 of 10, against 0 to 3 for docking) and loses on every NNRTI (3 to 4, against 7 to 8). The two metrics are measuring different things: distance orders the whole list well because DRMs are pocket-proximal, and then fills its top ten with residues that are close but not resistance mutations. Discriminating inside the pocket is what the energy contributes, and it contributes it only where the pocket is rigid.

Table 1: **Per-drug metrics.** DRM-recovery ROC-AUC and magnitude Spearman ρ vs. measured fold-resistance, per drug. n is the drug’s docked panel size, n_{DRM} its number of major DRMs.

target	drug	n	n_{DRM}	base rate	ROC-AUC	ρ
protease	ATV	258	33	0.128	0.533	−0.058
protease	DRV	209	30	0.144	0.462	0.033
protease	IDV	282	33	0.117	0.388	−0.085
protease	LPV	278	33	0.119	0.552	−0.013
protease	NFV	284	33	0.116	0.615	0.009
protease	SQV	280	33	0.118	0.672	0.138*
RT	EFV	423	30	0.071	0.738	0.078
RT	ETR	301	30	0.100	0.589	−0.003
RT	NVP	421	29	0.069	0.769	0.188**
RT	RPV	124	25	0.202	0.634	0.126

* $p < 0.05$, ** $p < 0.001$ for the Spearman correlation.

Table 2: **Per-drug precision@ k** , reported as (DRMs recovered)/ k . * marks $p < 0.05$ by an exact one-sided hypergeometric test against the drug’s own base rate.

drug	docking $\Delta\Delta G$			distance	base rate
	$k = 5$	$k = 10$	$k = 20$	$k = 10$	
<i>protease</i>					
ATV	3/5*	3/10	3/20	8/10	0.13
DRV	1/5	1/10	2/20	10/10	0.14
IDV	0/5	0/10	1/20	7/10	0.12
LPV	2/5	3/10	6/20*	8/10	0.12
NFV	2/5	2/10	4/20	7/10	0.12
SQV	1/5	2/10	5/20	7/10	0.12
<i>RT</i>					
EFV	4/5*	8/10*	15/20*	3/10	0.07
ETR	3/5*	7/10*	11/20*	3/10	0.10
NVP	4/5*	8/10*	15/20*	3/10	0.07
RPV	3/5	7/10*	9/20*	4/10	0.20

For completeness, Table 3 gives precision@ k on the *pooled* ranking (one ordering over all drug–mutation pairs of a target). The pooled protease view is the weaker of the two framings: only $k = 20$ clears $p < 0.05$, and pooled precision@10 is 2/10 against a 0.123 base rate. This is the same tension noted in §4.1, where protease shows tail enrichment without a usable ranking, and it is why we report the per-drug view as the primary one.

Table 4 gives every predictor with bootstrap 95% CIs; §4.1 quotes the point estimates from it. The intervals show that protease docking’s ROC-AUC contains 0.5 while RT’s excludes it, and that both HIVdb intervals sit far above every other predictor on both targets.

Table 3: **Pooled precision@ k** , one ranking over all pairs per target. $*p < 0.05$ (exact hypergeometric).

target	$k = 5$	$k = 10$	$k = 20$	base rate
protease	2/5	2/10	6/20*	0.123
RT	4/5*	8/10*	14/20*	0.090

Table 4: **Baselines with bootstrap 95% CIs** (5,000 resamples, seed 0). DRM-recovery ROC-AUC and magnitude Spearman ρ vs. measured fold-resistance. HIVdb’s ROC is partly circular against the expert DRM labels, so ρ is the fair column.

target	predictor	ROC-AUC [95% CI]	ρ [95% CI]
protease	docking $\Delta\Delta G$	0.510 [0.465, 0.557]	0.002 [−0.048, 0.052]
protease	mutation prevalence	0.719 [0.686, 0.751]	−0.078 [−0.127, −0.029]
protease	$ \Delta\text{volume} $	0.489 [0.445, 0.537]	0.142 [0.093, 0.191]
protease	HIVdb penalty	0.863 [0.828, 0.897]	0.360 [0.321, 0.395]
RT	docking $\Delta\Delta G$	0.695 [0.625, 0.768]	−0.106 [−0.164, −0.045]
RT	mutation prevalence	0.653 [0.608, 0.698]	0.127 [0.072, 0.182]
RT	$ \Delta\text{volume} $	0.653 [0.603, 0.701]	−0.009 [−0.065, 0.048]
RT	HIVdb penalty	0.929 [0.896, 0.959]	0.267 [0.209, 0.322]

Table 5 is the sharpest negative result in the paper: as a global ranker, the single number “how far is this residue from the ligand” outranks a full docking sweep on both targets, and on protease it also beats the expert system on the magnitude column. (At the top of the list the ordering reverses on RT; see Table 2.) Two readings are available and we do not think the evidence separates them. The deflationary one is that DRM status for a competitive inhibitor is close to a statement about pocket proximity, so a distance predictor is partly recovering the definition of the label; this is why we lead with the magnitude correlation against *measured* fold-resistance, where the same ordering holds but the margin is smaller. The constructive one is that the geometric localization of a mutation, which docking computes en route to its energy, is where the usable signal lives, and the scoring step discards rather than adds. The same conclusion arrives independently from the field-level ablation on the explanation side (§4.3), which is what makes us take it seriously.

The leave-one-drug-out jackknife quoted in §4.1 refits the pooled ROC-AUC with each drug held out in turn: protease stays within [0.50, 0.53] and RT within [0.67, 0.74], so neither the protease null nor the RT signal is carried by any single compound. Both are produced by `scripts/12_baselines.py`.

C Worked example: one mutation through the ablation

To make the attribution ablation concrete, we trace a single pair, indinavir (IDV) with the protease flap mutation **M46I**, through all three context conditions. M46I is the paradigm weak-prior case: an *accessory* mutation that does not touch the inhibitor.

Ground truth (IAS-USA 2022): “*Met46 is a flap residue that does not contact the inhibitor directly. M46I is an accessory flap mutation that alters flap dynamics and closure and compensates for active-site mutations, broadly contributing to PI resistance.*”

Table 6 shows the context each condition receives.

Full context (score 2). “*The M46I mutation increases hydrophobicity while maintaining nearly identical volume, subtly altering the hydrophobic landscape of the flap region 7 Å from indinavir without directly contacting it... reflects reduced entropic penalty for flap closure and weakened indirect stabilization of inhibitor binding geometry, characteristic of flap mutations that achieve resistance through conformational rather than direct contact mechanisms.*” The judge scored 2: the correct primary mechanism.

Table 5: **Geometry-only predictors vs. docking** on the same scored pairs, with bootstrap 95% CIs (scripts/15_geometry_baseline.py). “-min distance” ranks a mutation by how close its residue lies to the ligand; “contacts ligand” is the binary $< 4 \text{ \AA}$ flag. Neither requires a docking run.

target	predictor	ROC-AUC [95% CI]	magnitude ρ
protease	docking $\Delta\Delta G$	0.510 [0.466, 0.557]	0.002
protease	-min distance to ligand	0.963 [0.954, 0.971]	0.387
protease	contacts ligand (binary)	0.776 [0.741, 0.809]	0.134
protease	$ \Delta\text{volume} $	0.489 [0.442, 0.535]	0.142
RT	docking $\Delta\Delta G$	0.695 [0.622, 0.764]	-0.106
RT	-min distance to ligand	0.800 [0.738, 0.854]	0.130
RT	contacts ligand (binary)	0.660 [0.615, 0.705]	0.178
RT	$ \Delta\text{volume} $	0.653 [0.604, 0.702]	-0.009

Table 6: **Structural context supplied to the model** for IDV/M46I under each condition. *Corrupted* keeps the intrinsic chemistry (volume, hydrophobicity) but deranges the pipeline-derived geometry and energy across pairs.

field	full (real)	minimal	corrupted
min. distance to ligand ($\text{\AA}$)	7.0	—	4.1
distance to ligand centroid ($\text{\AA}$)	14.2	—	10.7
contacts ligand directly	False	—	True
region	flap	—	S2/S2' subpocket
Δvolume ($\text{\AA}^3$)	+3.8	—	+3.8
$\Delta\text{hydrophobicity}$	+2.6	—	+2.6
docking $\Delta\Delta G$ (kcal/mol)	-0.24	—	+0.66
judge score	2	0	0

306 **Minimal context** (score 0). “The M46I substitution replaces a larger, flexible methionine with a
307 smaller isoleucine, creating a structural cavity in the S1 subpocket where indinavir’s P1 ligand
308 group normally makes critical van der Waals contacts...” With only the drug and mutation name,
309 the model defaults to an active-site packing story and *contradicts* the documented mechanism.

310 **Corrupted context** (score 0). “... isoleucine’s shorter branched aliphatic chain cannot maintain
311 the same contact surface... the $\sim 4 \text{ \AA}$ minimum distance suggests the residue is in direct contact
312 with the ligand... destabilizes the inhibitor by approximately 0.66 kcal/mol, particularly affecting
313 the packing interactions critical for IDV’s potency in this region.”

314 The corrupted output is the informative one: the explanation *quotes the falsified values back* (the $4 \text{ \AA}$
315 distance, the S2/S2’ region, the 0.66 kcal/mol energy) and reaches the wrong mechanism because
316 of them. This is direct evidence that the generated mechanism is conditioned on the pipeline’s
317 structural fields rather than retrieved from the model’s prior on a famous mutation name, and it is
318 the single-instance version of the aggregate result in §4.3.

319 **Judge rationales.** The judge returns a one-sentence justification with every score; for this pair they
320 are, verbatim:

- 321 • *full* (2): “The explanation correctly identifies M46I as a flap mutation that operates through
322 conformational dynamics rather than direct contact, accurately describing how altered hy-
323 drophobic landscape and flap closure dynamics indirectly weaken inhibitor stabilization,
324 the core known mechanism.”
- 325 • *minimal* (0): “The explanation incorrectly claims M46I creates a structural cavity that dis-
326 rupts direct van der Waals contacts with indinavir’s P1 ligand group, contradicting the
327 known mechanism that Met46 is a flap residue that does NOT contact the inhibitor di-
328 rectly...”

329 • *corrupted* (0): “The explanation contradicts the known mechanism by claiming M46I re-
 330 duces binding through direct van der Waals contacts with indinavir, when the ground truth
 331 explicitly states Met46 is a flap residue that does not contact the inhibitor directly. . .”

332 Both failures turn on the same discriminating fact, whether residue 46 contacts the ligand, which is
 333 exactly the field the corruption flips (`contacts_ligand_directly`: False $\rightarrow$ True) and exactly
 334 the field group the field-level ablation identifies as most load-bearing (§4.3, *drop_distance*). The
 335 same pair appears in all four PIs that carry M46I (ATV, IDV, LPV, NFV) with the identical 2/0/0
 336 score pattern, so it is a reproducible instance rather than a single sampled outlier.

337 D Reproducibility

338 **Structures and docking.** Protease: PDB 3OXC, chains A/B, catalytic Asp25 dyad asym-
 339 metrically protonated, mutations applied to both chains of the homodimer; box centre
 340 (5.341, -1.893, 14.179), size 22³ Å. RT: PDB 3V81, p66/p51 heterodimer, NNRTI pocket; box
 341 centre (41.105, 52.332, 49.098) (occupancy-weighted centroid of the chain-A nevirapine ligand;
 342 note that 3V81 contains a second NVP copy whose inclusion moves the centroid into empty space),
 343 size 24³ Å; wildtype absolute affinities of -9 to -12 kcal/mol confirm the box is on the pocket.
 344 AutoDock Vina 1.2 (CPU) / Uni-Dock (GPU), exhaustiveness 16, 5 poses retained per run, best
 345 pose scored. Receptors cleaned with PDBFixer; mutants built by point mutagenesis with a frozen-
 346 environment clash-relief minimization of the mutated side chain only; ligands prepared with RDKit
 347 (AddHs, ETKDG) and meeko. Crystal waters, including the conserved flap water, are stripped.

348 **Docking-dropout sensitivity.** Pairs whose docking returned no valid pose are excluded, which
 349 would bias the benchmark if those pairs were DRM-enriched. They are not: **zero** of the 97 excluded
 350 RT pairs and zero of the 6 excluded protease pairs are major DRMs, against a 9.0% DRM rate
 351 among the 1,269 scored RT pairs (Fisher exact $p = 3.6 \times 10^{-4}$). All 30 RT and all 33 protease
 352 primary DRMs survive into the scored set. The residual worry is the reverse: the excluded pairs are
 353 non-DRMs that, had they docked, might have scored high and acted as false positives. We bound
 354 that directly by imputing every excluded pair at the *worst case* (a $\Delta\Delta G$ above every scored pair, i.e.
 355 ranked top and all wrong): RT’s pooled ROC-AUC falls only from 0.695 to **0.641**, still above chance,
 356 and protease is unchanged (0.510 to 0.508). Median imputation moves neither target. Produced by
 357 `scripts/15_geometry_baseline.py`.

358 **Scale, yield and cost.** Protease: 1,597 drug-mutation dockings attempted across 6 PIs, of which 6
 359 (0.4%) returned no valid pose, leaving **1,591 scored pairs over 285 mutant receptors**. RT: 1,366
 360 attempted across 4 NNRTIs, of which 97 (7.1%) returned no valid pose, leaving **1,269 scored pairs**
 361 **over 423 mutant receptors**. Pairs without a valid pose are excluded from every analysis rather than
 362 imputed; the RT failure rate is an order of magnitude higher than protease’s and is itself a caveat on
 363 the RT result, since the excluded set is not guaranteed to be missing at random. All 30 RT primary
 364 DRMs and all 33 protease primary DRMs are present in the scored set. The full RT sweep took
 365 19 minutes on a single A100 80GB with Uni-Dock. At roughly 2.7 s per receptor this is orders of
 366 magnitude cheaper per mutation than an alchemical free-energy protocol such as Hauser et al. [11]
 367 (which we did not benchmark directly), and that cost/accuracy trade is the one the paper is about.

368 **Judge robustness (detail for §4.2).** On a 21-item sample stratified for the judge’s hard (0/1)
 369 cases, GPT-5.6 agrees with Claude-Haiku 95% within one point but 48% exactly (κ 0.23/0.38
 370 unweighted/linear); Claude is systematically stricter. Re-grading all 138 protease ablation items
 371 with a second Claude model (Sonnet 5) reproduces the §4.3 contrasts: full-minimal +0.35 and
 372 full-corrupted +0.65, against Haiku’s +0.43 and +0.63 ($\kappa = 0.64/0.72$ unweighted/linear). The
 373 paired contrasts survive the change of judge; the absolute percentage does not.

374 **Models.** Generation and judging both use `claude-haiku-4-5` (`max_tokens` 1024 for generation,
 375 512 for judging); larger Claude models decline HIV-resistance *generation* but not judging, which
 376 is why the second judge (`claude-sonnet-5`) was available for the cross-model check. Scores
 377 are returned through a forced tool call against a fixed `report_score` schema, so every verdict is
 378 machine-readable and no free-text parsing is involved.

379 **Generation prompt** (verbatim system prompt; the enzyme noun and binding site are substituted per
 380 target, either “HIV-1 protease”/“active site” or “HIV-1 reverse transcriptase”/“allosteric NNRTI-
 381 binding pocket”): “You are a structural biologist specializing in HIV-1 protease drug resistance.

382 *Given structural context about a point mutation and its effect on inhibitor binding, provide a 2-*
 383 *3 sentence mechanistic hypothesis for WHY this mutation causes resistance to this specific drug.*
 384 *Ground your explanation in the structural facts provided (residue size/charge change, distance to the*
 385 *active site, subpocket). Be specific about molecular interactions (van der Waals contacts, hydrogen*
 386 *bonds, steric clashes, electrostatic changes). Do not speculate beyond the given data, and do not*
 387 *restate the numbers back — interpret them. No preamble.”*

388 The user message is templated as

```
389     Drug: {drug_full} ({drug}), a {drug_class}
390     Mutation: {mutation}
391     Predicted delta-delta-G: {ddg} kcal/mol (positive = weaker binding =
392     resistance)
393     Structural context:
394     {context JSON, indent 2}
395
396     Explain the structural mechanism by which {mutation} likely reduces
397     {drug_full} binding.
```

398 The *minimal* condition drops the $\Delta\Delta G$ line and the context JSON entirely, keeping the drug and
 399 mutation lines and the closing instruction; *full* and *corrupted* reproduce the production prompt
 400 verbatim (with `cite=False`), so the ablation’s control arm is the real pipeline rather than a re-
 401 implementation.

402 **Corruption procedure.** The corrupted condition replaces six pipeline-derived
 403 fields with those of another pair: `distance_from_ligand_centroid_angstrom`,
 404 `min_distance_to_ligand_angstrom`, `contacts_ligand_directly`, `region`,
 405 `delta_delta_g` and `delta_g` (the $\Delta\Delta G$ shown in the prompt follows the donor). Donors are
 406 assigned by a *derangement* (sampled with seed 0, rejecting any permutation with a fixed point), so
 407 no pair keeps its own values and every corrupted context is realistic-but-wrong. Intrinsic-chemistry
 408 fields (position, wildtype and mutant residue identity, and the volume/charge/hydrophobicity
 409 changes) are deliberately *not* corrupted, because the mutation string itself determines them and
 410 scrambling them would make the context self-contradictory rather than merely false. A consequence
 411 worth stating: on low-cardinality fields such as `region` a derangement can coincide with the true
 412 value by chance, so the corruption is a lower bound on the effect a directed (antipodal) corruption
 413 would produce.

414 **Contact-preserving corruption (control).** Because `contacts_ligand_directly` is close to the
 415 answer for a mechanism-classification question, we ran a stricter variant that isolates the rest of the
 416 geometry. Donors are drawn *within contact class*, so contacting residues donate only to contacting
 417 residues, and the flag is restored to its true value afterwards, so it is both correct and consistent with
 418 the donated distance (no self-contradictory prompt). The derangement is over distinct *mutations*
 419 rather than (drug, mutation) pairs, since the same mutation recurs across drugs with identical geom-
 420 etry; the contact flag is preserved on all 46 protease pairs and the distance genuinely changes on 41
 421 of them.

422 Faithfulness under this condition is 1.630 (72% correct), against 1.152 (41%) for the full derange-
 423 ment and 1.783 (80%) for real context. Three paired Wilcoxon contrasts follow: contact-preserved
 424 beats full corruption by $\Delta = +0.478$ ($p = 5.8 \times 10^{-4}$), while real context still beats contact-
 425 preserved by $\Delta = +0.152$ ($p = 0.035$; 9 up, 3 down, 34 tied). The decomposition is therefore
 426 about 76% of the corruption effect attributable to the contact boolean and about 24% to distance,
 427 subpocket and energy. We report this because it materially qualifies §4.3: the pipeline does con-
 428 tribute beyond a single flag, but the headline *corrupted* arm overstates that contribution by roughly
 429 a factor of four. Produced by `scripts/17_contact_preserving_corruption.py`.

430 **Field-level ablation groups.** Each condition strips one group from the context: *drop_ΔΔG*
 431 removes `delta_delta_g` and `delta_g` and nulls the $\Delta\Delta G$ prompt line; *drop_distance* re-
 432 moves `distance_from_ligand_centroid_angstrom`, `min_distance_to_ligand_angstrom`
 433 and `contacts_ligand_directly`; *drop_subpocket* removes `region`; *drop_chem* removes
 434 `volume_change`, `volume_change_A3`, `charge_change` and `hydrophobicity_change`. All
 435 other fields and the prompt are unchanged.

436 **Judge prompt** (verbatim system prompt): “*You are evaluating whether an AI-generated explana-*
 437 *tion of HIV-1 protease drug resistance correctly identifies the known structural mechanism. Score 0*”

438 *if the explanation contradicts the known mechanism, 1 if it is consistent but vague or misses the key*
 439 *interaction, 2 if it correctly identifies the primary structural basis. Judge only structural/mechanistic*
 440 *agreement, not writing style, and give a one-sentence justification.”* The enzyme noun is substituted
 441 per target. The judge sees the mutation, the drug, the ground-truth mechanism and the explanation.
 442 It never sees the condition label, the $\Delta\Delta G$, or which system produced the explanation.

443 **Statistics.** Permutation tests use 5,000 label shuffles and bootstrap CIs 5,000 resamples, seed 0. En-
 444 richment ranks by $\Delta\Delta G$ descending. Paired ablation contrasts use the one-sided Wilcoxon signed-
 445 rank test at the (drug, mutation) level, which is the appropriate paired non-parametric test for the
 446 bounded, tied, non-Gaussian 0/1/2 score differences we observe; a paired t -test would assume a
 447 normality the difference distribution violates. Precision@ k significance is the exact hypergeometric
 448 test described in Appendix B. We report many tests without multiplicity correction; the headline
 449 contrasts survive at $p < 10^{-3}$.

450 **Ground truth.** Protease mechanisms are hand-curated from the IAS-USA 2022 mutation list and
 451 the primary structural literature, with the source recorded per entry. RT mechanisms (22 entries)
 452 are built by a tool-using agent that queries live PubMed for each mutation, cites the papers it read,
 453 and self-verifies that the stated mechanism is supported before the entry is admitted; entries failing
 454 that check are dropped rather than included unverified. Because all RT ground truth is agent-built
 455 while all protease ground truth is human-curated, absolute faithfulness percentages are not strictly
 456 comparable *across* targets. The within-target ablation contrasts, which hold ground truth fixed, are.

457 **Software.** Python 3.11; pandas 2.2.3, numpy 1.26.4, scipy 1.15.3, scikit-learn 1.7.0, anthropic
 458 0.74.0; RDKit, OpenMM, PDBFixer, AutoDock Vina and meeko via conda (environment.yml).
 459 The pipeline is reproduced by scripts/01–14 in order; each number in this paper is produced by
 460 a named script, and the mapping is in the repository README.

461 **Data availability.** All docking results, explanations under every ablation condition, ground-truth
 462 mechanisms, judge scores with justifications, and the derived tables in this appendix are released
 463 with the code, so every figure and table here can be regenerated without re-running docking or
 464 re-querying any model.

465 E Failure-mode analysis

466 The method makes one implicit assumption: *that resistance is a local, contact-mediated pertur-*
 467 *bation of ligand binding, visible in a single rigid structure.* Every failure we observe violates one
 468 clause of that sentence, and which clause it violates predicts the failure. Table 7 organises them; the
 469 point is that these are not scattered noise but four identifiable regimes, three of which are properties
 470 of the *pipeline*, not of the model.

471 Two consequences worth stating plainly. First, eight of the nine faithfulness misses are attributable
 472 to what the pipeline could not represent rather than to the language model; only A98G is a plain
 473 model error. The E138 cluster is the clearest case: the LLM was faithful to its input and the input
 474 was mechanistically misleading. This is the sharpest illustration of the paper’s own thesis, since
 475 geometric context is what carries the explanation, so wrong or incomplete geometry is what breaks
 476 it. It also names a concrete fix (an inter-subunit contact feature) rather than a vague one. Second,
 477 the protease and RT failures are complementary: protease fails because resistance is non-local, RT’s
 478 residual failures because the locality *measure* is incomplete. Neither is repaired by a better language
 479 model.

480 This analysis is also the one part of the faithfulness evaluation that does not depend on an LLM.
 481 Each diagnosis above is checkable against artifacts outside the loop: that E138 sits 43.0 Å from the
 482 chain-A ligand is a fact of the structure file, and that Met46 and Leu90 are non-contact residues is
 483 recorded in the IAS-USA list. It is a partial check only, validating *why* the misses occur, not the
 484 absolute faithfulness percentage, but it does establish that the score-0 cases are real failures with
 485 identifiable causes rather than judge artifacts.

Table 7: **Failure modes**, grouped by the assumption each violates. Faithfulness misses are every score-0 explanation on both targets (3/46 protease, 6/55 RT).

case	assumption violated	mechanism of the failure
Protease active site	locality (resistance is a co-evolving network)	Substrate-envelope-constrained, multi-mutation resistance [13, 16]. Pooled ROC-AUC 0.51; indinavir is <i>below</i> chance at 0.388 with 0/10 precision@10. Rigid single-structure pocket affinity has no representation of the network.
L90M (all 3 protease score-0 misses: ATV, IDV, SQV)	contact-mediated	L90M does not contact the inhibitor; its documented mechanism is indirect distortion of active-site geometry. The model instead invokes dimer-interface destabilization or direct hydrophobic-contact loss. All three misses are the <i>same mutation</i> across three drugs.
E138A/G/K (4 of 6 RT score-0 misses)	contact-mediated & single-subunit	The true mechanism is the E138-K101 salt bridge across the p51/p66 interface [8]. The pipeline reports <code>min_distance_to_ligand</code> = 43.0 Å (measured to the chain-A ligand) and <code>contacts_ligand_directly</code> = False, so the model reasons <i>correctly</i> from a numerically true fact to the wrong conclusion, “~43 Å away, negligible impact.” The context has no inter-subunit coupling feature.
A98G (RT)	none (model error)	The context is correct and sufficient; the model simply dis-misses a real accessory contribution as “essentially neutral.” The only miss of the nine attributable to the model rather than the pipeline.
L100I (RT)	single rigid pose	The model attributes resistance to a change of binding mode; a single-pose rigid protocol cannot represent one.

F An LLM-free structural-consistency check

Every check in §4.2 is internal to a family of language models. This appendix adds one that is not, by exploiting the fact that a single claim in each explanation is mechanically verifiable against the receptor: *does the mutated residue touch the ligand?*

Method. For each of the 303 explanations (46 protease and 55 RT pairs $\times$ 3 conditions) we take the true minimum residue–ligand distance from the structure (contact iff < 4 Å) and, separately, classify what the explanation *asserts* using a fixed regex lexicon in which negation is resolved first, so that “without directly contacting it” is scored as a non-contact claim rather than a contact one. Explanations making no explicit claim *abstain* and are excluded from accuracy rather than counted wrong. No language model participates in the measurement; the lexicon and the per-explanation calls are released with the code (`scripts/16_structural_consistency.py`).

Coverage. 139 of 303 explanations (46%) make an explicit contact claim. Abstention is itself condition-dependent: stripped of the distance field, the *minimal* arm makes far fewer checkable claims ($n = 14$), which is why that cell is the weakest of the three.

Result A: the ablation reproduces without an LLM (Table 8). Explanations written with the real structural context get the contact question right 88% of the time; stripped of it, 57%; given deranged geometry, 55% (Fisher exact $p = 0.013$ full vs. minimal, $p = 5.7 \times 10^{-5}$ full vs. corrupted). The ordering, the magnitude and the near-equality of the minimal and corrupted arms all match the judge-scored result of §4.3, obtained by an entirely different route.

Result B: the judge tracks an external criterion. Explanations that are structurally *wrong* by this test receive a mean judge score of 1.02, versus 1.51 for structurally correct ones (Mann–Whitney one-sided $p = 6.1 \times 10^{-4}$; $n = 42$ and 97). The LLM judge is therefore not merely agreeing with other LLMs: its scores move with a criterion computed from a PDB file.

What this does and does not establish. It is a check on one axis, over the 46% of explanations that commit to it, using a lexicon we wrote; it is not a substitute for expert grading of the full mechanistic claim, and we do not report it as one. What it does establish is that the two results most exposed to the circularity objection, the ablation contrast and the judge’s discrimination, both survive a measurement taken outside the LLM family.

Table 8: **Contact-claim accuracy by condition**, computed from receptor geometry with no language model. Accuracy is over explanations making an explicit claim; abstentions are excluded.

condition	n claims	accuracy	protease / RT
full (real context)	60	0.883	0.862 / 0.903
minimal (no context)	14	0.571	0.500 / 0.667
corrupted (deranged)	65	0.554	0.583 / 0.537

G Extended related work

Supervised ML on genotype–phenotype data. geno2pheno learns decision trees and SVMs mapping sequence to phenotypic resistance [3]; deep architectures (MLP, bidirectional RNN, CNN) have been trained across 18 antiretrovirals [20]; large-scale gradient boosting and random forests over clinical sequence corpora recover both known and novel resistance mutations [4]; and protein language-model embeddings push mean AUC to ≈ 0.97 [10]. Each needs a labeled resistance dataset for the drug in question, which is the requirement docking does not carry.

Alternatives to rigid docking. Other structure-aware approaches include molecular field mapping over homology models [18] and, at the accurate-but-expensive end, alchemical free-energy calculations (88% accuracy, ≈ 1.1 kcal/mol RMSE on 144 clinical Abl mutations [11]) at orders of magnitude more compute per mutation than the 2.7 s per receptor reported above. Learned co-folding models are a fast-moving alternative [5], though benchmarking shows deep docking can produce physically invalid poses and generalize poorly off-distribution [7].

Perturb-and-rejudge. Perturbing the input and re-scoring with an LLM judge is an established family: Porbeck et al. [17] vary debugging context and judge the resulting failure explanations, and Afolabi et al. [1] use causal ablation to test whether stated medical reasoning drives the answer. Our design is not new; its transfer to *structural* input is, as is the verification of Appendix F, since one claim in every explanation here is checkable against a PDB file with no language model. LLMs are increasingly deployed as tool-using agents in chemistry [6], but the faithfulness of their mechanistic claims in this domain is rarely measured.

H What each baseline needs to exist

Table 4 shows HIVdb beating docking on both targets. The argument for running docking anyway rests on *what data each predictor requires*, which Table 9 makes explicit rather than asserting.

Table 9: **Information each predictor requires** for a given (compound, target). “Novel compound” asks whether the predictor can be computed for a scaffold with no clinical history.

predictor	receptor structure	clinical isolates	expert curation	measured phenotypes	novel compound?
docking $\Delta\Delta G$	✓	—	—	—	yes
$ \Delta \text{ side-chain volume} $	—	—	—	—	yes
mutation prevalence	—	✓	—	—	no
Stanford HIVdb penalty	—	✓	✓	✓	no
supervised ML [3, 20, 4, 10]	—	✓	—	✓	no

Prevalence needs a corpus of clinical isolates genotyped *and* tested against that drug; HIVdb needs, in addition, expert curation of per-mutation per-drug penalties distilled from accumulated genotype–phenotype and outcome data; supervised models need labelled phenotypes for the drug they predict. None of these exist for a compound that has never been dosed in a patient. Docking needs only a receptor structure and the candidate’s chemistry.

What “prospective” does and does not mean here. Our $\Delta\Delta G$ comes from a physics/empirical docking engine (AutoDock Vina, or Uni-Dock’s GPU reimplementations of the same scoring function), not a learned one. This matters for how the result should be evaluated. Deep docking and co-folding models [5] carry billions of parameters trained on experimental complexes, so a held-out test set is mandatory to show they learned structural biology rather than memorizing training proteins,

547 which is exactly the generalization gap Buttenschoen et al. [7] document. Vina is not parameter-free
548 either: its handful of scoring-term weights were fit once by regression on general protein–ligand
549 crystal data. But that fit involved no resistance phenotype, no HIVdb record and nothing in this
550 evaluation. The precise claim is therefore not “nothing is fit” but *nothing in the $\Delta\Delta G$ path is fit*
551 *to resistance data*, which is what rules out leakage into this benchmark. Consequently there is no
552 resistance-specific training set to hold out, and the per-drug numbers in Table 1 are not a proxy for
553 prospective performance. For a new compound against these targets, they *are* it.

554 Two boundaries on that. First, the *mutation panel* is target-derived from clinical data: docking
555 tells you how a compound fares against a mutation, not which mutations to test. So the method
556 is prospective per compound but not yet per target, and validating a genuinely novel target needs
557 a resistance dataset that, by construction, a novel target lacks. Second, the pipeline’s one heavily
558 trained component is the LLM, for which the memorisation concern is real and live; the *minimal*
559 condition of §4.3 is precisely its control, measuring what the model produces from the drug and
560 mutation name alone.

561 **The honest qualification.** Being training-free is not by itself the virtue: $|\Delta\text{volume}|$ is also com-
562 putable for any novel compound, needing nothing but the mutation string, and on RT it reaches
563 ROC-AUC 0.653 [0.603, 0.701] against docking’s 0.695 [0.625, 0.768], intervals that overlap sub-
564 stantially (Table 4). On protease it is 0.489, at chance, as docking is. And the ligand-distance
565 baseline of Table 5, also structure-only and also training-free, outranks docking globally on both
566 targets. So the defensible claim is narrower than “docking is the only prospective option”: among
567 prospectively-available predictors, docking is the best we measured only at the top of the ranking
568 and only on RT (Table 2), and on the global metric it is not separated from a zero-cost physicochem-
569 ical proxy. Separating them, and testing transfer to a new target, are the experiments that would
570 complete the argument.